

# The Hidden Common Ground in Business Agents and Agentic Coding

## O'Reilly API Superstream: APIs and the Agentic Shift

**Erik Wilde**  
Head of Enterprise Strategy, Jentic

July 14, 2026



This work is licensed under a CC  
Attribution 4.0 International License

# Erik Wilde

- Head of Enterprise Strategy at Jentic
- Ambassador for the OpenAPI Initiative (OAI)
- [LinkedIn](#) [linkedin.com/in/erikwilde](https://www.linkedin.com/in/erikwilde)
- [YouTube](#) [youtube.com/ErikWilde](https://www.youtube.com/ErikWilde)



# Outline

1. [Agents & Platforms](#) [3]
2. [Disconnect: AI-First without API-First](#) [5]
3. [Disconnect: Symbolic Security](#) [4]
4. [Conclusions](#) [3]



# Coding Agents and Business Agents

- *Coding Agents*

- They change *how the business does things*
- Tool side speed-up
- Getting existing work done faster

- *Business Agents*

- They change *what the business does*
- Business-side reinvention
- From narrow off-the-shelf to broad and ambitious



# The Platform Principle

- You don't need a platform to ship *one* thing
- Platforms are for shipping *many* similar things at scale
  - Shared substrate beats many bespoke stacks
  - Capabilities, governance, security → built for reuse
- That's where both kinds of agents are heading
  - Coding agents: already at scale; platform spirit emerging
  - Business agents: scale is the next problem to solve



# What We Found

- Many reports on AI scaling but none on *this* disconnect
- We surveyed both sides
  - IT / Platform engineering
  - Business
  - 110+ responses across 50 organizations
  - 16 long-form interviews
- Found 6 disconnects
- Today: 2 of them
  - [Disconnect: AI-First without API-First](#)
  - [Disconnect: Symbolic Security](#)



# Outline

1. [Agents & Platforms](#) [3]
2. [Disconnect: AI-First without API-First](#) [5]
3. [Disconnect: Symbolic Security](#) [4]
4. [Conclusions](#) [3]



# AI-First without API-First: The Problem


- Agents getting connected to APIs (directly or via MCP)
- API quality and AI-readiness not checked first
- Broken or sub-optimal workflows now exposed to agents
- Old problems carry into the new AI layer
- API problems may surface as AI problems



# AI-First without API-First: The Data

While dev teams are mostly responsible for identifying how to connect business systems to coding agents, Almost half (45%) of respondents are connecting coding agents to their APIs, representing a slightly higher rate of integration than amongst business agents, where 2 in 5 respondents (41%) are leveraging internal APIs.

**Are your APIs used for business agents?**



...But only 1 in 5 respondents (21%) noted have processes in place to assess the readiness of their APIs for AI consumption.

**Have your APIs been assessed for AI readiness?**



# AI-First without API-First: The Consequences

- Platform foundations skipped under C-suite AI pressure
- AI pilots: isolated results, scaling not part of the plan
- Bad processes get scaled up, not fixed




# Tactics: MCP-Washing is not a Strategy

- Make AI-readiness a gate before agent integrations
- Reuse and refactor existing API governance
- Fix bad APIs & workflows *before* agents touch them
- Don't paper over old problems with MCP



# Next Step: AI-Readiness Scorecard

- APIs are AI enablement
- Assess where you are today
- Six dimensions of AI-readiness
- Web: [jentic.com/scorecard](https://jentic.com/scorecard)
- CLI: `npx @jentic/api-scorecard-cli score openapi.yaml`



# Outline

1. [Agents & Platforms](#) [3]
2. [Disconnect: AI-First without API-First](#) [5]
3. [Disconnect: Symbolic Security](#) [4]
4. [Conclusions](#) [3]



# Symbolic Security: The Problem

- Agents are less predictable than traditional code
- Coding agents produce code: reviewable before it runs
- Business agents produce actions: live as they're decided
- Both need limits but controls in place are mostly symbolic



# Symbolic Security: The Data

We know that agents should be able to be tracked individually to ensure that individual agent deployments and risks can be understood.

**55%** of survey respondents give agents their own API keys so that usage can be tracked.

**47%** noted that there is a predefined model for delegating user rights to business agents.

**47%** survey respondents use downscoping to limit agent usage.

But even fewer are actually monitoring the agent use, so these safeguards are not providing any security for new agentic systems:

**47%** of survey respondents did not know if agent usage was being monitored for data and security breaches or other concerns



# Symbolic Security: The Consequences

- Some controls exist, but no one is watching
- Vulnerabilities and unusual patterns invisible
- Context and action creep: agents access more than they need
- Optionality takes a hit when experiments feel unsafe



# Symbolic Security: Tactics

- Add a control layer between agents and APIs
- Default-deny: grant only what's needed
- Monitor what agents do, not just *that* they have keys
- Alerts + thresholds → human-in-the-loop



# Outline

1. [Agents & Platforms](#) [3]
2. [Disconnect: AI-First without API-First](#) [5]
3. [Disconnect: Symbolic Security](#) [4]
4. [Conclusions](#) [3]



# AI Platform Tactics

- One agent platform is the ideal, not the overnight plan
- Treat both kinds of agents as one demand stream
- Share governance, security, observability
- Spot parallel platforms and manage strategically



# More Disconnects in the Report



*This is not an O'Reilly form. Your information will be submitted directly to Jentic and processed under their responsibility and privacy notice.*



# Thank You!

-  [YouTube\\_youtube.com/ErikWilde](https://youtube.com/ErikWilde)
- [Linkedin\\_linkedin.com/in/erikwilde](https://linkedin.com/in/erikwilde)

